

A Collection of Thoughts on the Keccak/SHA-3 Construction

Alexander Scheel

Iowa State University

In this paper, we discuss new techniques for the cryptanalysis of Keccak/SHA-3, though none of the resulting attacks scale to full versions of the hash function. We show an explicit construction for the inverses of the core round functions, the permutation orders of the round functions, and an attack which exploits low-order cycles. Further, we introduce new logical cryptanalysis techniques and show the existence of partial fixed points to produce 2-1 block collisions.

MATH 490 | Advisor: Dr. Clifford Bergman | SHA-3 | [hash_framework](#)

1. A Series of Introductions

This paper presents the work of the author towards the completion of the requirements for the MATH 490 Independent Study course under Dr. Clifford Bergman at Iowa State University of Science and Technology. This work was an extension of the author's Honors Project under Dr. Eric Davis and utilized the resulting [hash_framework](#) project (viewable under the [cipherboy](#) account on GitHub.com). Additional artifacts related to this project can be seen in the [keccak-attacks](#) repository. A permanent location for this document is in the [papers](#) repository.

Within this section are a series of introductions which provide necessary background on the topics of cryptographic hash functions, the development of Keccak/SHA-3, and its structure. While certain sections can be skipped if the reader has the prerequisite knowledge, hopefully all readers find the material engaging and useful. Following these introductions, this paper presents the analysis of the Keccak/SHA-3 hash function before concluding with a final evaluation and further work.

The remaining sections outside of the introductions discuss 1) various mathematical properties of the core round functions, 2) exhaustive collision searches on small instances, 3) correlation matrices, and 4) fixed point attacks.

Introduction on the Topics of Cryptographic Hash Functions. While there are many applications of pure mathematics, few are as demanding and shrouded in secrecy as cryptography. Cryptography exists because of the fundamental need of civilizations, governments, and individuals to keep secrets secure from devoted adversaries. While modern cryptography combines the disciplines of mathematics, computer science, and computer engineering, prior to the turn of the 20th century, cryptography lacked much of its modern rigor.

Within cryptography's collection of algorithms, few are as useful as hash functions have proven to be to cryptographers and non-cryptographers alike. A hash function maps arbitrary length binary strings to binary strings of a fixed length; typically we denote such a function h . Because the inputs are of arbitrary size, by the pigeonhole principle, there must exist at least one collision: two blocks

$x \neq y$ such that $h(x) = h(y)$. However, to be considered cryptographically secure, finding such a collision must be hard (on the order of $2^{\frac{b}{2}}$ for b the number of bits in the output).

Further, a cryptographic hash function must be resistant finding a preimage: for a given output value v , finding an input x such that $h(x) = v$ should be roughly on the order of 2^b . Given such an attack, it is possible to find a collision: choose a random input, compute its hash, then use the preimage attack to find a different input for that hash; these two inputs form a collision pair. Thus a preimage attack is stronger than a collision attack.

However, if the attacker is allowed to choose only one message (that is, the challenger chooses the other), this forms a second preimage attack. More formally, for a given fixed message x , the attacker seeks a $y \neq x$ such that $h(x) = h(y)$. While a collision attack is only valid for some such messages, x , a preimage attack has to be valid for all such x ; thus it is a stronger attack than collision attacks, but weaker than a preimage attack as it presupposes an existing message.

Introduction on the Usages of Hash Functions. Studying the integrity of hash function is important because hash functions are a fundamental building block of both cryptographic and non-cryptographic systems. While a body of literature developed around attacking the previous hash functions (of MD4, MD5, SHA-1 and SHA-2), many of these attacks do not directly translate to Keccak/SHA-3 because its construction differs substantially. Further, many of these attacks started only as generic attacks with low utility, and only with substantial work, were they translated to highly specialized attacks on specific use cases. Below is a discussion of various places where cryptographic hash functions are used, and what types of attacks they are most susceptible to.

The most common use case for hash functions are in signatures. Often times the cryptography underlying asymmetric signatures (where one party can sign a message and any party can verify that the signature and message is authentic) can only sign fixed length messages, often only in the thousands of bytes. However, longer data frequently must be signed. Thus, most implementations of signature schemes use a cryptographically secure hash function and sign the digest instead of the actual message. This can be seen in PGP, TLS, and PKI. Often, these implementations are most vulnerable to a second preimage attack, as one message is already provided. However, when the attacker has control over both messages, a collision attack will suffice.

Another common use case for hash functions are in message integrity verification. In this case, a hash of the message is sent along side of the message to ensure that the message hasn't been tampered with. While often this signature should be signed, it is frequently used in cases where the authenticity of the hash isn't questioned. This includes verifying contents of files on disk or in version control systems such as Git. In both of these cases, second preimage and collision attacks are the most useful for an attacker.

Introduction on the Scope of the Independent Study. This independent study focused on a single hash function, Keccak, which was chosen by NIST to be the next standardized algorithm, SHA-3 [1]. This hash function differs from previously standardized algorithms in two major ways: one, it is not based on the Merkle-Damgård construct and is rather based on the Sponge construction, and two, the internal round function is not a compression function and is instead a permutation. Outside of a few theorems done by hand and verifying the lack of low-order correlations, we focused on using automated tools to learn more about the structure of SHA-3.

In the late 1990s, F. Massacci started analyzing cryptographic constructions using Boolean Satisfiability [2]; however, applying this to hash functions is a largely understudied field. E. Homsirikamol tried brute forcing collisions in SHA-3 to limited results [3], and P. Morawiecki performed similar work for preimages [4]. However, neither contributed useful, general cryptanalysis techniques to scale attacks against reduced rounds (and reduced sizes) to larger numbers of rounds. Our work begins

exploring new techniques and several consequences in the following sections.

Introduction to Terminology. This section contains a collection terminology useful for discussing SHA-3 and Boolean Satisfiability.

We define $\Sigma = \{0, 1\}$ to be the alphabet.

We define the set $W = \{1, 2, 4, 8, 16, 32, 64\}$ to be the powers of two typically used in constructing the Keccak widths; $w = 64$ is standardized for use in SHA-3, though any power of two can be used.

For a given $w \in W$, we define $S_w = \Sigma^{25w}$, to be the set of all binary strings of length $25w$; these represent the possible states (which are binary strings that map onto an indexable array A described later) and each round permutation maps $S_w \mapsto S_w$.

A binary string $s \in S_w$ can be indexed directly (in a zero-indexed manner, i.e., the starting bit of s is denoted $s[0]$), or via a 3 dimensional structure of size 5 by 5 by w . This is typically done in conjunction with an uppercase letter, $A = s$, and then $A[x, y, z] = s[w(5y + x) + z]$. This is in accordance with FIPS 202 [1].

Introduction on the Development of Keccak/SHA-3. NIST standardized Keccak as SHA-3 in FIPS 202 [1]; Keccak was chosen as the winner of the corresponding hash function design competition started in 2009. Designed by the Keccak Team (G. Bertoni, J. Daemen, M. Peeters, G. Van Assche, and R. Van Keer [5]), Keccak features a novel construction.

Keccak is built around the concept of a state cube: a 5 by 5 by w cube of bits (where w is a power of 2). By increasing w , a higher security margin can be given, at the expense of increasing the computational intensity. However, because internal round operations on this state cube are bijective, it can be implemented efficiently in hardware with minimal gate and propagation delays. Typically this state is called A when represented as a cube, and s when represented as a (flat) bit string. There are five major functions which compose to form a single round: θ , ρ , π , χ , and ι ; each maps bijectively over the entire state space (i.e., elements of 2^{25w}).

θ is an eleven way XOR and is thus a linear map. For all $0 \leq x \leq 4$ and $0 \leq z < w$, define:

$$C[x, z] := \bigoplus_{0 \leq y \leq 4} A[x, y, z]$$

$$D[x, z] := C[(x - 1) \bmod 5, z] \oplus C[(x + 1) \bmod 5, (z - 1) \bmod w]$$

Then:

$$A'[x, y, z] := A[x, y, z] \oplus D[x, z]$$

ρ is a pure permutation function, permuting the locations of bits in a lane (the z component) of the state cube. For all $0 \leq z < w$: $A'[0, 0, z] := A[0, 0, z]$. Then let $(x, y) = (1, 0)$, and for $0 \leq t \leq 23$:

$$\forall 0 \leq z < w : A'[x, y, z] := A[x, y, (z - \frac{(t+1)(t+2)}{2}) \bmod w]$$

$$\text{let } (x, y) = (y, (2x + 3y) \bmod 5)$$

π is another pure permutation function, permuting locations along the face of the cube. For all $0 \leq x \leq 4, 0 \leq y \leq 4, 0 \leq z < w$:

$$A'[x, y, z] := A[(x + 3y) \bmod 5, x, z].$$

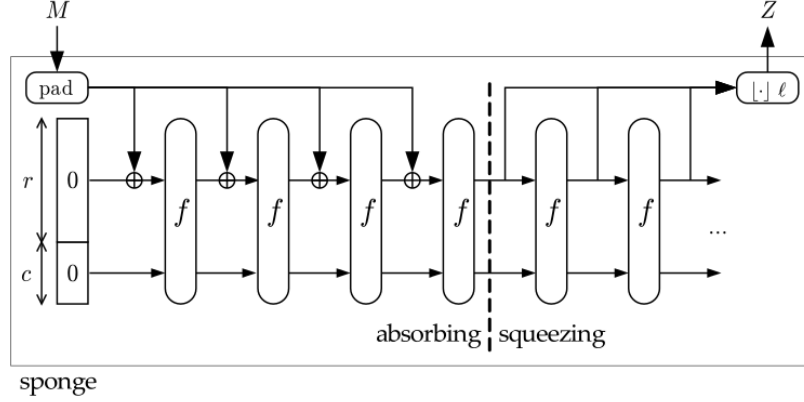


Fig. 1. Graphical description of the sponge construction. Courtesy of the [Keccak Team](#).

χ is the only non-linear function and is of degree two (with inverse of degree three). For all $0 \leq x \leq 4, 0 \leq y \leq 4, 0 \leq z < w$:

$$A'[x, y, z] := A[x, y, z] \oplus (\neg A[(x+1) \bmod 5, y, z] \wedge A[(x+2) \bmod 5, y, z]).$$

ι is a fixed-value XOR useful for preventing trivial fixed points, and is the only function dependent on the current round number. Refer to page 16 of FIPS 202 ([1]) for its specification; it is not discussed extensively in this paper. It takes an additional parameter, c , the current round number.

The a single round of the core round function function, r is defined as follows:

$$r(A, c) = \iota(\chi(\pi(\rho(\theta(A)))), c)$$

And the entire core round function, $\text{KECCAK-}f$, is thus:

$$\text{KECCAK-}f(A) = r(r(\dots r(A, 0), \dots), 22), 23)$$

The sponge construction can be described as follows. Fix w as a power of two in W . Then define a state, s , as above to be a binary string of length $25w$. Fix a security margin m (usually 228, 256, 384, or 512 bits in the SHA-3 standard). Define $i_m = 25w - 2m$ be the input margin, and $o_m = 2m$ be the output margin. For a message M , assume it is padded to a multiple of i_m bits according to the scheme in FIPS 202 (`pad10*1`). Start with initial state $s_0 = 0^{25w}$. For each block M_i , XOR M_i into the first $|M_i|$ bits of state s_i and define:

$$s_{i+1} = \text{KECCAK-}f(s_i \oplus M_i)$$

Then if s_k is the final state, take the first m bits ($s_k[0..m]$) as the output digest. Additionally if more than $2m$ bits are required, take the first $2m$ bits, then compute $s_{k+1} = \text{KECCAK-}f(s_k)$ and continue taking at most the first $2m$ bits from s_{k+1} , repeating as necessary. This is demonstrated in Figure 1.

To have a security margin of m bits, the output needs to have m bits; however, if the attacker has complete control of all $25w$ bits of state, they would be able to choose an output with prefix m easily because Keccak is bijective with a known, computable inverse. Thus, by restricting the initial state to have a suffix of $2m$ bits of zeros, the attacker must not only restrict the output to have the correct prefix, but also the input to have the correct suffix. In this way, the authors claim

that Keccak/SHA-3 has collision resistance equal to preimage resistance; finding a collision is of equivalent difficulty as finding a preimage, and thus Keccak/SHA-3 is more secure than necessary [6].

Further, SHA-3 is secure from trivial multiple-block collisions due to the restriction of allowing i_m -bits of input state to be changed directly via a block. Thus, for a multiple-block collision attack to be possible, the remaining $25w - i_m = 2m$ bits must be the same between the internal state of a collision pair, allowing two different blocks b and b' to cancel their initial prefix differences. All of these attacks are constrained input/output problems and thus reduce to SAT; thus, considerable work is required to successfully perform an attack.

2. Mathematical Properties of the Five Round Functions

In the following section, we detail various mathematical properties of the five permutation functions which make up the core round function of Keccak. In most cases, we seek to provide mathematical proofs of these properties. In all cases, we rely on external code and Boolean Satisfiability for computerized proofs where proofs are not directly provided. We do note the introduction of a cleaner form of χ^{-1} than in existing literature.

Of θ . In this section, we show that θ is bijective, that XOR distributes through θ , give a method for finding the inverse of θ , and give the order of θ .

Let w be a fixed power of 2. Since S_w is of finite size, it suffices to show that, $\forall x, y \in S_w$, $x \neq y \Rightarrow \theta(x) \neq \theta(y)$. Assume the hypothesis: then $x \oplus y \neq 0^{25w}$.

Lemma 2.1.

$$\theta(a) = 0^{25w} \iff a = 0^{25w}$$

Proof. First, note that $\theta(0^{25w}) = 0^{25w}$ due to construction of θ having an even number of XORs. Suppose that there existed some other input, x , to θ such that $\theta(x) = 0^{25w}$. Note that, x with odd number of 1s (that is, $|\{i : x_i = 1\}| \equiv 1 \pmod{2}$), $\theta(x)$ likewise has an odd number of ones due to θ having an even number of XORs. Similarly, for any x with has an even but non-zero number of ones, then there exists index (i, j, k) such that $\theta(x)[i, j, k] = 1$. This is because, for any pair of indices (i, j, k) and (i', j', k') such that $x[i, j, k] = 1 = x[i', j', k']$, there exists an index (a, b, c) such that $\theta(x)[a, b, c]$ is dependent on only one of (i, j, k) or (i', j', k') , by construction of θ . Thus the kernel of θ is 0^{25w} . □

Lemma 2.2. $\forall a, b \in S_w$,

$$\theta(a \oplus b) = \theta(a) \oplus \theta(b)$$

Proof. This follows from the definition of θ : note that θ is composed entirely of XORs and that XOR is commutative and associative. For any index (i, j, k) , we want to show that

$$\theta(a \oplus b)[i, j, k] = \theta(a)[i, j, k] \oplus \theta(b)[i, j, k]$$

Note that:

$$\begin{aligned}
\theta(a \oplus b)[i, j, k] &= (a[i, j, k] \oplus b[i, j, k]) \\
&\oplus \left(\bigoplus_{0 \leq y \leq 4} (a[(i-1) \bmod 5, j, k] \oplus b[(i-1) \bmod 5, j, k]) \right) \\
&\oplus \left(\bigoplus_{0 \leq y \leq 4} (a[(i+1) \bmod 5, j, (k-1) \bmod w] \oplus b[(i+1) \bmod 5, j, (k-1) \bmod w]) \right) \\
&= (a[i, j, k] \oplus \left(\bigoplus_{0 \leq y \leq 4} a[(i-1) \bmod 5, j, k] \right) \oplus \left(\bigoplus_{0 \leq y \leq 4} a[(i+1) \bmod 5, j, (k-1) \bmod w] \right)) \\
&\oplus (b[i, j, k] \oplus \left(\bigoplus_{0 \leq y \leq 4} b[(i-1) \bmod 5, j, k] \right) \oplus \left(\bigoplus_{0 \leq y \leq 4} b[(i+1) \bmod 5, j, (k-1) \bmod w] \right)) \\
&= \theta(a)[i, j, k] \oplus \theta(b)[i, j, k]
\end{aligned}$$

□

Combining Lemma 2.1 and Lemma 2.2, we have that:

$$\begin{aligned}
x \oplus y = 0^{25w} &\iff \theta(x \oplus y) = \theta(0^{25w}) \\
&\iff \theta(x \oplus y) = 0^{25w} \\
&\iff \theta(x) \oplus \theta(y) = 0^{25w}
\end{aligned}$$

and hence θ is bijective.

To construct the inverse of θ , note that since θ is a linear map, there exists a matrix M such that $\theta(x) = Mx$, when x is treated as a column vector of bits. M is described by a square matrix of size $25w$ by $25w$, where for a given row i :

$$M_{i,j} = 1 \iff \theta(x)[i] \text{ depends on } x[j]$$

Since we know that θ is invertible, M must be invertible and thus M^{-1} describes the inverse of θ .

Lastly, we have computed the order of the permutation θ for all w . We reproduce them here without proof; they were found by randomized search, and verified with SAT for $w = 1, 2, 4$ and 8 . In general, the order is given by the expression $3 \times w$.

w	Order
1	3
2	6
4	12
8	24
16	48
32	96
64	192

Of ρ . Since ρ is a simple permutation of the location of bits, it holds that $\rho(a \oplus b) = \rho(a) \oplus \rho(b)$.

It is obvious that the order of the ρ permutation is w : this follows from the construction of ρ .

Of π . Since π is a simple permutation of the location of bits, it holds trivially that $\pi(a \oplus b) = \pi(a) \oplus \pi(b)$.

It is obvious that the order of the π permutation is 24: this follows from the construction of π .

Of χ . The order of the χ function is 4, and is independent of the width w . This is because χ is independent of both y and z coordinate. We will show algebraically that the order of χ is 4, however, it has also been verified with SAT. We first present three basic lemmas without proof:

Lemma 2.3.

$$\neg(a \oplus b) = \neg a \oplus b$$

Lemma 2.4.

$$(a \oplus b) \wedge (c \oplus d) = (a \wedge c) \oplus (a \wedge d) \oplus (b \wedge c) \oplus (b \wedge d)$$

Lemma 2.5.

$$\neg(a \oplus b) \wedge (c \oplus d) = (\neg a \oplus b) \wedge (c \oplus d)$$

We define a row of the state cube, A to have elements a_1, a_2, a_3, a_4 , and a_5 . After applying χ to this row, we define the resulting output row to be b_1, b_2, b_3, b_4, b_5 , where:

$$\begin{aligned} b_1 &= a_1 \oplus (\neg a_2 \wedge a_3) \\ b_2 &= a_2 \oplus (\neg a_3 \wedge a_4) \\ b_3 &= a_3 \oplus (\neg a_4 \wedge a_5) \\ b_4 &= a_4 \oplus (\neg a_5 \wedge a_1) \\ b_5 &= a_5 \oplus (\neg a_1 \wedge a_2) \end{aligned}$$

And similarly for $c_1 \dots c_5, d_1 \dots d_5, e_1 \dots e_5$. Thus, to show χ has order 4, it suffices to show that $e_1 = a_1$, (since χ has rotational symmetry; if $e_1 = a_1$, then $e_2 = a_2 \dots e_5 = a_5$).

Lemma 2.6.

$$\begin{aligned} c_1 &= b_1 \oplus (\neg b_2 \wedge b_3) \\ &= (a_1 \oplus (\neg a_2 \wedge a_3)) \oplus (\neg(a_2 \oplus (\neg a_3 \wedge a_4)) \wedge (a_3 \oplus (\neg a_4 \wedge a_5))) \\ &= a_1 \oplus (\neg a_2 \wedge a_3) \oplus (\neg a_2 \wedge a_3) \oplus (\neg a_2 \wedge \neg a_4 \wedge a_5) \\ &\quad \oplus (a_3 \wedge \neg a_3 \wedge a_4) \oplus (a_3 \wedge a_4 \wedge \neg a_4 \wedge a_5) \text{ (by Lemma 2.5)} \\ &= a_1 \oplus (\neg a_2 \wedge \neg a_4 \wedge a_5) \end{aligned}$$

And:

$$\begin{aligned} c_2 &= a_2 \oplus (\neg a_3 \wedge \neg a_5 \wedge a_1) \\ c_3 &= a_3 \oplus (\neg a_4 \wedge \neg a_1 \wedge a_2) \\ c_4 &= a_4 \oplus (\neg a_5 \wedge \neg a_2 \wedge a_3) \\ c_5 &= a_5 \oplus (\neg a_1 \wedge \neg a_3 \wedge a_4) \end{aligned}$$

Thus, to show $e_1 = a_1$:

$$\begin{aligned}
e_1 &= d_1 \oplus (\neg d_2 \wedge d_3) \\
&= c_1 \oplus (\neg c_2 \wedge \neg c_4 \wedge c_5) \\
&= (a_1 \oplus (\neg a_2 \wedge \neg a_4 \wedge a_5)) \oplus (\neg(a_2 \oplus (\neg a_3 \wedge \neg a_5 \wedge a_1)) \\
&\quad \wedge \neg(a_4 \oplus (\neg a_5 \wedge \neg a_2 \wedge a_3))) \\
&\quad \wedge (a_5 \oplus (\neg a_1 \wedge \neg a_3 \wedge a_4))) \\
&= a_1
\end{aligned}$$

Hence χ has order 4 and $\chi^{-1} = \chi^3$. The Keccak reference states: “We refer to [20, Section 6.6.2] for an algorithm for computing the inverse of χ .” [6]; this points a reference in J. Daemen’s thesis [7]. However, given χ , χ^3 is easy to compute and thus χ^{-1} is also easy to compute. Alternatively, using the above, it is easy to see that:

$$d_1 = a_1 \oplus (\neg a_2 \wedge a_3) \oplus (\neg a_2 \wedge \neg a_4 \wedge a_5)$$

Lastly, note that XOR does not distribute over χ due to the introduction of the and.

Of ι . Note that since ι is an XOR with a fixed value, it is obvious that ι is a bijection: for any w , for any i , and for all $x \in S_w$, $\iota(\iota(x, i), i) = x$, since $x \oplus \iota_i \oplus \iota_i = x$. Hence, ι is its own inverse and hence ι is bijective since the inverse is well defined for all $x \in S_w$.

Lastly, note that it is trivial that the order of the ι permutation is 2 by construction (due to the XOR).

Evaluation of the Orders of Composition of Permutations. In this section, we discuss how the above five permutation functions compose, and the orders of the resulting compositions. We limit our discussion to $w = 1$ for the interests of exhaustive search: 2^{25} is possible on commodity hardware, 2^{50} would require additional resources. In general, we find that these permutations interact non-trivially, resulting in cycles of mixed size. We reproduce these results in the table below:

Function	Order	Fixed Points	Number of Cycles	List of Cycles
θ	3	2097152	2	1, 3
ρ	1	33554432	1	1
π	24	4	8	1, 2, 3, 4, 6, 8, 12, 24
χ	4	32	3	1, 2, 4
$\chi \circ \pi$	17360392635484575518934418947500880 $\approx 2^{113.741}$	3	28	Not Reproduced
$\pi \circ \rho \circ \theta$	24	4	8	1, 2, 3, 4, 6, 8, 12, 24
$\chi \circ \pi \circ \rho \circ \theta$	418144575651966378899040573720 $\approx 2^{98.399}$	3	27	Not Reproduced
r_1	320185339723133697023127516600 $\approx 2^{98.014}$	0	14	Not Reproduced
r_2	$\approx 2^{130.726}$	0	12	Not Reproduced
r_3	$\approx 2^{164.242}$	1	16	Not Reproduced
r_4	$\approx 2^{211.609}$	0	16	Not Reproduced
r_5	$\approx 2^{131.878}$	2	16	Not Reproduced
r_6	$\approx 2^{190.743}$	3	18	Not Reproduced
r_7	$\approx 2^{218.017}$	0	18	Not Reproduced
r_8	$\approx 2^{190.137}$	0	18	Not Reproduced
r_9	$\approx 2^{188.483}$	0	14	Not Reproduced
r_{10}	$\approx 2^{154.432}$	0	19	Not Reproduced
r_{11}	$\approx 2^{223.948}$	1	18	Not Reproduced
r_{12}	$\approx 2^{108.579}$	0	12	Not Reproduced
r_{13}	$\approx 2^{209.785}$	1	18	Not Reproduced
r_{14}	$\approx 2^{170.621}$	0	14	Not Reproduced
r_{15}	$\approx 2^{196.507}$	2	22	Not Reproduced
r_{16}	$\approx 2^{172.643}$	1	16	Not Reproduced
r_{17}	$\approx 2^{221.160}$	1	22	Not Reproduced
r_{18}	$\approx 2^{203.510}$	0	20	Not Reproduced
r_{19}	$\approx 2^{250.748}$	3	20	Not Reproduced
r_{20}	$\approx 2^{183.937}$	1	19	Not Reproduced
r_{21}	$\approx 2^{158.852}$	2	15	Not Reproduced
r_{22}	$\approx 2^{111.874}$	1	12	Not Reproduced
r_{23}	$\approx 2^{230.807}$	0	24	Not Reproduced
r_{24}	$\approx 2^{140.355}$	1	14	Not Reproduced

Note that, for $w = 1$, ρ is the identity function and is thus omitted from the tables above. Note that r_n denotes the first n rounds of SHA-3.

While on first glance, large orders make it appear that the hash function is strong, SHA-3 also defines an XOF (or eXtensible Output Function) construct. This allows SHA-3 to act as a pseudo-random number generator: after hashing an input, the internal state of SHA-3 has reached some value S . Suppose we want to request m bits of state. If k exceeds our security margin $k/2$, we take $k/2$ bits, permute $S \rightarrow S'$ according to KECCAK- f , and repeat until all m bits have been retrieved.

An ideal permutation function for an XOF would ensure that there exists one random cycle through all possible states (and thus contain a cycle of 2^{25} , in this case). However, the order of the r_n rounds of SHA-3 far exceed 2^{25} and all of them have several cycles. This suggests that the actual security margin of the XOF construct is far less than the theoretical value. However, this analysis needs to be extended to at least $w = 4$ (performing 2^{100} iterations of the hash function core which is not feasible) to fully verify this.

3. Exhaustive Collision Searches

By modeling the problem with SAT, we were able to recreate the work of prior authors ([3], [4]). However, like previous authors, our results do not scale to useful results. For $w = 1$, we have exhaustively searched the collision space; the results of this search are reproduced below:

w	r	Number of collisions
1	1	1024
1	2	502
1	3	543
1	4	525
1	5	488
1	6	518
1	7	532
1	8	498
1	9	506
1	10	522
1	11	506
1	12	503
1	13	495
1	14	522
1	15	485
1	16	540
1	17	467
1	18	490
2	1	$> 2^{20}$
2	2	$> 2^{14}$
2	3	> 74

Note that the results for $w = 2$ are incomplete; the search was terminated after one week. Further, all of these used an effective margin of 512-bits. In our limited time, we were unable to find any useful patterns in this data.

4. Correlation Matrices

Another area of study was the resistance of Keccak to correlation attacks. That is, given some structured input (in this case, a valid block where the remaining bits are all zeros), does SHA-3 emit any useful two bit correlations across rounds? Here, we constructed a program to generate matrices of correlation for all two-variable boolean functions given some state. We then performed Monte-Carlo simulation; for $w = 8$ and $r = 24$, we found no useful correlations under sufficiently large starting states (> 16 million). Thus, Keccak is secure from cross-round first and second order correlations; all distinguisher attacks must thus rely on at least third-order correlations. For $w = 8$, checking all third-order correlations across 24 rounds becomes computationally infeasible and thus was not performed, since it is not likely to yield useful results.

5. Fixed Point Attacks

In this section, we introduce two potential attacks against Keccak/SHA-3 using fixed points in the underlying round function. However, neither of these attacks have been proven possible with current analysis except for small values of w and small numbers of rounds. The first is a possible attack using full fixed points in the core round function, while the latter describes a category of partial fixed points.

Full Fixed Points. One theoretical attack against SHA-3 is by using a fixed point. If there existed a state x such that $h(x) = x$, for h the round function, then for all additional blocks, $h(x|b) = h(x \oplus b)$. While fixed points can and do occur (see 24-rounds in Table 2), using a fixed-point attack is more subtle in practice: it is unlikely that a fixed point is a valid block (that is, x contains a suffix that is 0^m for the current margin, m). Thus one must find a series of blocks b_1, b_2 such that $h(b_1) \oplus b_2 = x$;

this in turn extends the attack from being $h(x|b) = h(x \oplus b)$ to:

$$\begin{aligned} x &= h(h(b_1) \oplus b_2) \Rightarrow \\ h(x|b) &= h(b_1|b_2|b) = h(b_1|h_2 \oplus b) \\ &= h(b_1|h_2|0^{25w}|b) \end{aligned}$$

However finding b_1, b_2 amounts to a preimage attack and is thus unlikely to occur.

For $w = 1, 2$, we have verified that no fixed points occur which are valid blocks for $1 \leq r \leq 24$, where r is the number of rounds, at margins of 4 and 8 bits.

Partial Fixed Points. Another theoretical attack is using a partial fixed point. Suppose there exists a block x such that $h(x) = y$ is also a block (that is, y has a suffix of 0^m). In this case, the following is a valid collision, for all blocks b :

$$h(x \oplus b) = h(y|b)$$

We have verified the existence of these partial fixed points for small w and number of rounds. Furthermore, they are more readily found using SAT than finding a fixed point is.

For $w = 1$, we reproduce a table of quantities of partial fixed points per round below for the first 8 rounds:

Rounds	Quantity	Input	Output
1	512	FFFFFFFFFFFFFFFFFFFFFFFF	FFFFFFFFFFFFFFFFFFFFFFFF
2	567	FTFTTTTFTTTTFTTFTTTTT	TTTFTFTTFTTTFTTFTTTTT
3	527	FTTFTTTTFTTTTTTTTTTTTT	TTTFTTTTFTTFTTFTTTTT
4	545	TTTFTTTTFTTFTTTTTTTTT	FFTFTTTTFTTFTTFTTTTT
5	488	FFFFFFFFTFTTTTTTTTTTTTT	TFTTFTTTTTTFTTFTTTTT
6	533	FTTFTTTTFTTTTFTTTTTTTTT	TTTFTTTTFTTTTFTTTTT
7	510	TFTFTTTTFTTTTFTTTTTTTTT	FTTTTTTTTTTFTTTTFTTTTT

Note that the above table is only for an effective margin of 256-bits (4 bits when $w = 1$). For an effective margin of 512-bits (8 bits), we found no such partial fixed points.

However, brute forcing a partial fixed point for larger values of w and r quickly grows impossible. A similar technique to the aforementioned differential matrices could extend these techniques to larger values of w and r . Instead of building a model to check for a specific number of differences, build a model to check for a specific number of zeros in the inputs and outputs. This could provide a reduction in number of search paths for the SAT solver, and allow for more explicit parallelism.

6. Conclusions and Further Work

Overall, Keccak/SHA-3 remains secure to many attacks from SAT solvers. The majority of the techniques discussed here do not scale to the full $w = 64$, rendering them ineffective for attacking SHA-3 in general. More efficient techniques are thus necessary for finding collisions in SHA-3; this likely involves finding ways in which structure in small values of w translate to structure in larger values of w .

Future work includes scaling the partial fixed point attacks via intermediate models, evaluating weaknesses in alternate constructions, and improving the run time of marginal and differential analysis. However, without significant breakthroughs in logical cryptanalysis, most of these attacks are unlikely. Further, it is not immediately obvious what effect low permutation order has on the possibility of collisions; introducing permutation-based cryptanalysis could result in new techniques applicable to SHA-3. Additional work could include automatic lemmaizing of found constraints and including them in future searches.

7. Bibliography

1. (2015) Fips pub 202, sha-3 standard: Permutation-based hash and extendable-output functions. U.S.Department of Commerce/National Institute of Standards and Technology. <http://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.202.pdf>.
2. Massacci F (year?) Logical cryptanalysis. <http://disi.unitn.it/~massacci/CryptoSAT/>.
3. Homsirikamol E, Morawiecki P, Rogawski M, Srebrny M (2012) *Security Margin Evaluation of SHA-3 Contest Finalists through SAT-Based Attacks*, eds. Cortesi A, Chaki N, Saeed K, Wierzchoń S. (Springer Berlin Heidelberg, Berlin, Heidelberg), pp. 56–67. <https://eprint.iacr.org/2012/421.pdf>.
4. Morawiecki P, Srebrny M (2010) A sat-based preimage analysis of reduced keccak hash functions (Cryptology ePrint Archive, Report 2010/285). <https://eprint.iacr.org/2010/285.pdf>.
5. Team K (2008 - 2017) Team keccak | home (online). <https://keccak.team/>.
6. Bertoni G, Daemen J, Peeters M, Van Assche G, Van Keer R (2011) The keccak reference - version 3.0. <https://keccak.team/files/Keccak-reference-3.0.pdf>.
7. Daemen J (1995) Cipher and hash function design strategies based on linear and differential cryptanalysis. http://jda.noekeon.org/JDA_Thesis_1995.pdf.